

1. Faraway 2.1. The dataset `teengamb` concerns a study of teenage gambling in Britain. Fit a regression model with the expenditure on gambling as the response and the sex, status, income and verbal score as predictors. Present the output.

Call:

```
lm(formula = gamble ~ sex + status + income + verbal)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-51.082	-11.320	-1.451	9.452	94.252

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	22.55565	17.19680	1.312	0.1968
sex	-22.11833	8.21111	-2.694	0.0101 *
status	0.05223	0.28111	0.186	0.8535
income	4.96198	1.02539	4.839	1.79e-05 ***
verbal	-2.95949	2.17215	-1.362	0.1803

Signif. codes: 0 *** 0.1 standard error: 22.69 on 42 degrees of freedom

Multiple R-squared: 0.5267, Adjusted R-squared: 0.4816

F-statistic: 11.69 on 4 and 42 DF, p-value: 1.815e-06

(a) What percentage of variation in the response is explained by these predictors?

This is given by R^2 , i.e. 53%.

(b) Which observation has the largest (positive) residual? Give the case number.

The 39th indexed response has the minimum residual of -51.082 .

(c) Compute the mean and median of the residuals.

Theoretically we expect that the mean residual is 0. When calculated it, is $< 3 \times 10^{-16}$. The median residual is -1.451 .

(d) Compute the correlation of the residuals with the fitted values.

Again, theoretically we expect this to be 0, and, when calculated, it is $< 3 \times 10^{-17}$.

(e) Compute the correlation of the residuals with the income.

This quantity was computed to be less than 6×10^{-17} . So long as the model assumptions are met, this is expected to be 0.

(f) For all other predictors held constant, what would be the difference in predicted expenditure on gambling for the male compared to a female?

This is given by the coefficient corresponding to `sex`. Since 0 was coded for males, we expect that males gamble 22.12 £ more a week than females on average when all other explanatory variables are held constant.

2. Faraway 2.3. In this question, we investigate the relative merits of methods for computing the coefficients. Generate some artificial data by:

```
> x <- 1:20
> y <- x+rnorm(20)
```

Fit a polynomial in x for predicting y . Compute $\hat{\beta}$ in two ways – by `lm()` and by using the direct calculation described in the chapter. At what degree of polynomial does the direct calculation method fail? (Note the need for the `I()` function in fitting the polynomial, that is, `lm(y ~ x + I(x^2))`).

Solution:

Rather than following Faraway's suggestion, we employ the function `vandermonde.matrix` in the `matrixcalc` library. E.g. a design matrix for a quadratic model on $x = c(1,2,3,4,5)$ is given by

```
> vandermonde.matrix(1:5,3)
      [,1] [,2] [,3]
[1,]     1     1     1
[2,]     1     2     4
[3,]     1     3     9
[4,]     1     4    16
[5,]     1     5    25
```

We can now easily loop through the possible design matrices given by `A = vandermonde.matrix(x,20)`. The results follow:

```
> for (i in 2:20) { print(solve(crossprod(A[,1:i]),A[,1:i]),crossprod(A[,1:i],y))) }
      [,1]
[1,] -0.2054614
[2,]  1.0541753
      [,1]
[1,] -0.564601934
[2,]  1.152122725
[3,] -0.004664162
```

... % OUTPUT OMITTED

```
      [,1]
[1,] -8.717609e-01
[2,]  1.602359e+00
[3,] -1.733619e-01
[4,]  2.380631e-02
[5,] -1.394025e-03
[6,]  2.864698e-05
Error in solve.default(crossprod(A[, 1:i], A[, 1:i]), crossprod(A[, 1:i], :
  system is computationally singular: reciprocal condition number = 3.54243e-18
```

Hence, solving the normal equations is only numerically possible for up to fifth degree polynomials, and the design matrix will necessarily have two columns that

are “numerically linearly dependent” in the sense that $A^T A$ will be numerically singular.

Using `lm(...)` seems to use a more stable method as a full list of parameters are reported up to 13th degree polynomials.

```
> A = A[,-1] % remove intercept column
> for (i in 1:19) { # This will use lm( ... ) to find the parameters
+   out = lm(y ~ A[,1:i])
+   print(sprintf("----- i = %d -----",i))
+   print(summary(out))
+ }
```

... % OUTPUT OMITTED

```
[1] "----- i = 13 -----"
```

Call:

```
lm(formula = y ~ A[, 1:i])
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.53293	-0.21569	0.02872	0.24767	1.43972

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.899e+01	1.009e+02	-0.188	0.856
A[, 1:i]1	3.506e+01	2.576e+02	0.136	0.896
A[, 1:i]2	-1.816e+01	2.598e+02	-0.070	0.946
A[, 1:i]3	8.977e-01	1.409e+02	0.006	0.995
A[, 1:i]4	2.516e+00	4.672e+01	0.054	0.959
A[, 1:i]5	-1.098e+00	1.012e+01	-0.108	0.917
A[, 1:i]6	2.332e-01	1.489e+00	0.157	0.880
A[, 1:i]7	-2.996e-02	1.511e-01	-0.198	0.849
A[, 1:i]8	2.473e-03	1.059e-02	0.233	0.822
A[, 1:i]9	-1.322e-04	5.032e-04	-0.263	0.800
A[, 1:i]10	4.434e-06	1.546e-05	0.287	0.782
A[, 1:i]11	-8.481e-08	2.769e-07	-0.306	0.768
A[, 1:i]12	7.063e-10	2.196e-09	0.322	0.757
A[, 1:i]13	NA	NA	NA	NA

... % OUTPUT OMITTED

□


3. Faraway 2.4. The dataset `prostate` comes from a study on 97 men with prostate cancer who were due to receive a radical prostatectomy. Fit a model with `lpsa` as the response and `lcavol` as the predictor. Record the residual standard error and the R^2 . Now add `lweight`, `svi`, `lpbh`, `age`, `lcp`, `pgg45` and `gleason` to the model one at a time. For each model record the residual standard error and the R^2 . Plot the trends in these two statistics.

Solution:

For `lpsa ~ lcavol`, the calculated R^2 is 0.5394 and the residual standard error is 58.915.

In the following table we add `lweight`, `svi`, `lpbh`, `age`, `lcp`, `pgg45` and `gleason` to the model in order and record the resulting R^2 and RSS.

Explanatory Variables	R^2	RSS
<code>lcavol</code>	0.539	58.91
<code>...+lweight</code>	0.586	52.97
<code>...+svi</code>	0.626	47.78
<code>...+lpbh</code>	0.637	46.48
<code>...+age</code>	0.644	45.53
<code>...+lcp</code>	0.645	45.40
<code>...+pgg45</code>	0.654	44.20
<code>...+gleason</code>	0.655	44.16




Note that as explanatory variables are added, the model predictive power (R^2) increases and that the distance from the predictions to the data (the residual sums of squares) decreases as expected. The trends are non-linear, and indicate “diminishing returns” in the sense that adding an additional parameters reduces the RSS and increases R^2 less and less.

□

4. Faraway 2.5. Using the `prostate` data, plot `lpsa` against `lcavol`. Fit the regression of `lpsa` on `lcavol` and `lcavol` on `lpsa`. Display both regression lines on the plot. At what point do the two lines intersect?

Solution:

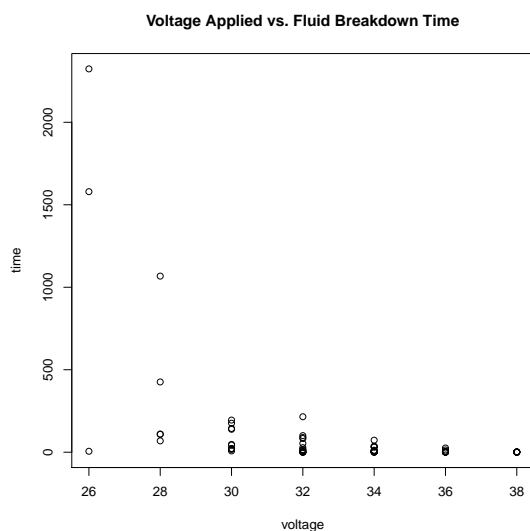


Note that the least squares line for `lcavol ~ lpsa` minimizes the vertical distance from the lines to the points, and the least squares line for `lpsa ~ lcavol` minimizes the horizontal distance. Hence, their intersection occurs at the respective means, i.e. at $(\overline{\text{lpsa}}, \overline{\text{lcavol}}) = (2.478, 1.350)$

□


5. In an industrial laboratory, under uniform conditions, batches of electrical insulating fluid were subjected to constant voltages until the insulating property of the fluids broke down. Seven different voltage levels, spaced 2 kilovolts (kV) apart from 26 to 38 kV, were studied. The measured responses were the time, in minutes, until breakdown. These data can be found on the course webpage under the file `fluid.txt`. [Data from W.B. Nelson, Schenectady, NY: GE Co. Tech. Report 71-C-011 (1970).]

(a) Construct a scatterplot of breakdown time vs. voltage and describe the association. Can you suggest a nonlinear model that could be fit to these data? IF you were to fit such a model, would the homogeneous error variance assumption be valid?



I might suggest a reciprocal model in this case, i.e. the response time y_i is modeled by the reciprocal voltage $1/x_i$. This would require that a fluid have an arbitrarily long breakdown time for no applied voltage. Note that any curve that fits this model will *not* have homogeneous additive error.

(b) As mentioned in class, transformation on variables are performed for many reasons. In a regression setting, in addition to linearizing the relationship between two variables, transformations sometimes will correct violations of other model assumptions, such as the error variance homogeneity assumption. Take the (natural) log of the breakdown times, plot these against the voltages, and describe the resulting association between the variables. Discuss whether this transformation also homogenized the error variance.



(c) Fit the simple linear regression model with $\log(\text{breakdown time})$ as the response

and voltage as the explanatory variable, and report the regression equation. What breakdown time is predicted for a voltage of 35 kV?

(d) Give an interpretation of the slope in terms of the voltages and the raw breakdown times. In other words, don't interpret the slope in terms of the log breakdown times.

(e) Compute and interpret the coefficient of determination R^2 .

(f) Construct a residual plot and normal quantile plot on the residuals for this model and explain what the plots indicate about the regression model.

(g) Suppose in addition to these two variables, a third variable (let's just call it temperature, even though I completely made it up!) is measured in this experiment, as given in the 3rd column of the data set. Fit a linear regression model with $\log(\text{breakdown time})$ as the response and voltage & temperature as explanatory variables. Report the ANOVA table for this regression fit and explain in a few sentences what is happening with this model.

6. The 1994 World Almanac reports heights and numbers of stories for notable tall buildings in North America. A random sample of 60 buildings among those in the list for which dates a completion were available was taken. The data appear under the filename `building.txt` on the course web page in four columns. The columns in order consist of the building name, year of completion, height (feet), and number of stories. Analyze the data and write a brief statistical report including a summary of statistical findings and any relevant graphical displays, where you address the following questions. What is the distribution of the number of stories as a function of height? What is the mean height per story? Are there any buildings that have particularly fewer stories than expected for their height? Are there any buildings that have particularly more stories than expected for their height? Is there any indication that the number of stories per height has changed over the time period represented here?

7. For the simple linear regression model ($y_i = \beta_0 + \beta_1 x_i + \epsilon_i$), show that $\text{Cov}(\bar{y}, \hat{\beta}_1) = 0$ whenever $\text{Var}(\epsilon_i) = \sigma^2$ for all i and $\text{Cov}(\epsilon_i, \epsilon_j) = 0$ for all $i \neq j$.

Solution:

Using a result from calculating the variance of β_1 in the notes and bilinear properties of covariance, we have

$$\begin{aligned} \text{cov}(\bar{y}, \hat{\beta}_1) &= \text{cov}\left(\frac{1}{n} \sum_{i=1}^n y_i, \frac{1}{s_{xx}} \sum_{j=1}^n (x_j - \bar{x}) y_j\right) \\ &= \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^n (x_j - \bar{x}) \text{cov}(y_i, y_j) \\ &= \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^n (x_j - \bar{x}) \text{cov}(\epsilon_i, \epsilon_j) && \text{since } y_i = \beta_0 + \beta_1 x_i + \epsilon_i \\ &= \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x}) \sigma^2 = 0. \end{aligned}$$

□

8. Prove the last equality on page 51 of the class notes for the MSE, namely that:

$$\text{MSE} = \frac{1}{n-2} \left[(n-1)s_y^2 - \hat{\beta}_1^2 (n-1)s_x^2 \right].$$

Solution:

Recall that $\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$ and $\hat{\beta}_1 = \frac{s_{xy}}{s_{xx}}$. Continuing from page 51,

$$\begin{aligned}
 MSE &= \frac{1}{n-2} \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2 \\
 &= \frac{1}{n-2} \sum_{i=1}^n \left((y_i - \bar{y}) + \hat{\beta}_1 (\bar{x} - x_i) \right)^2 \\
 &= \frac{1}{n-2} \sum_{i=1}^n \left((y_i - \bar{y})^2 - 2\hat{\beta}_1 (y_i - \bar{y})(x_i - \bar{x}) + \hat{\beta}_1^2 (x_i - \bar{x})^2 \right) \\
 &= \frac{1}{n-2} \left((n-1)s_y^2 - 2\frac{s_{xy}}{s_{xx}}s_{xy} + \left(\frac{s_{xy}}{s_{xx}} \right)^2 s_{xx} \right) \\
 &= \frac{1}{n-2} \left((n-1)s_y^2 - \hat{\beta}_1(n-1)s_x^2 \right).
 \end{aligned}$$

□